

STAARS Seminar
Grenoble, June 18th

A Reconfigurable Component Model for HPC

Vincent Lanore¹, Christian Pérez²

¹ : ENS de Lyon, ² : Inria
LIP laboratory, Avalon Team
France

High-Performance Computing

Goal: run the biggest possible applications

- from astrophysics, meteorology, industry...
- up to millions of years of sequential computing time

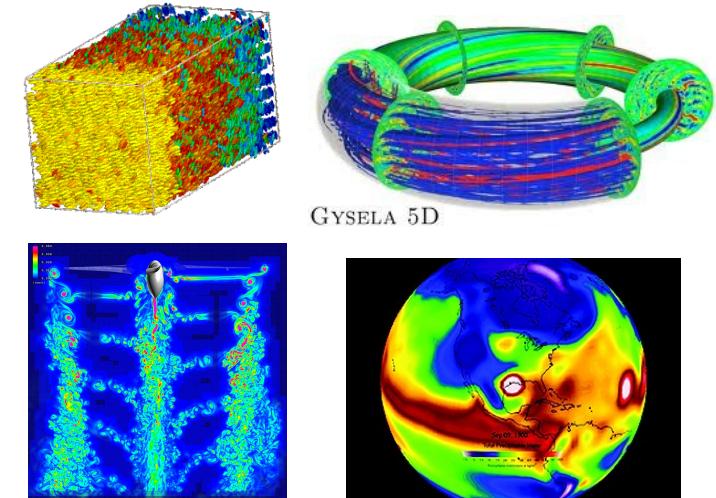
using cutting edge hardware

- very parallel

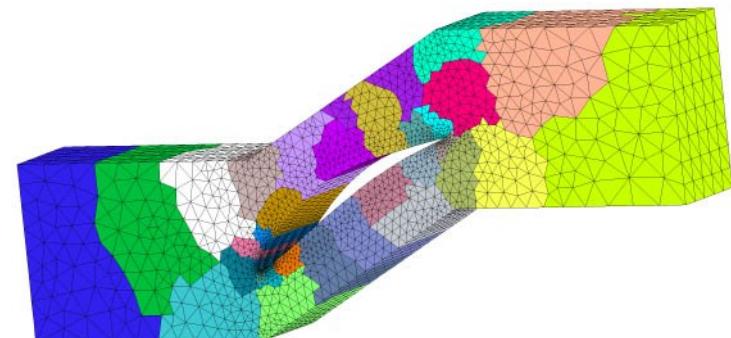
Challenge: scalability



Tianhe-2
3,120,000 cores
source: top500.org



Various application domains



A scientific mesh-based simulation
source: NTUA, school of mechanical engineering

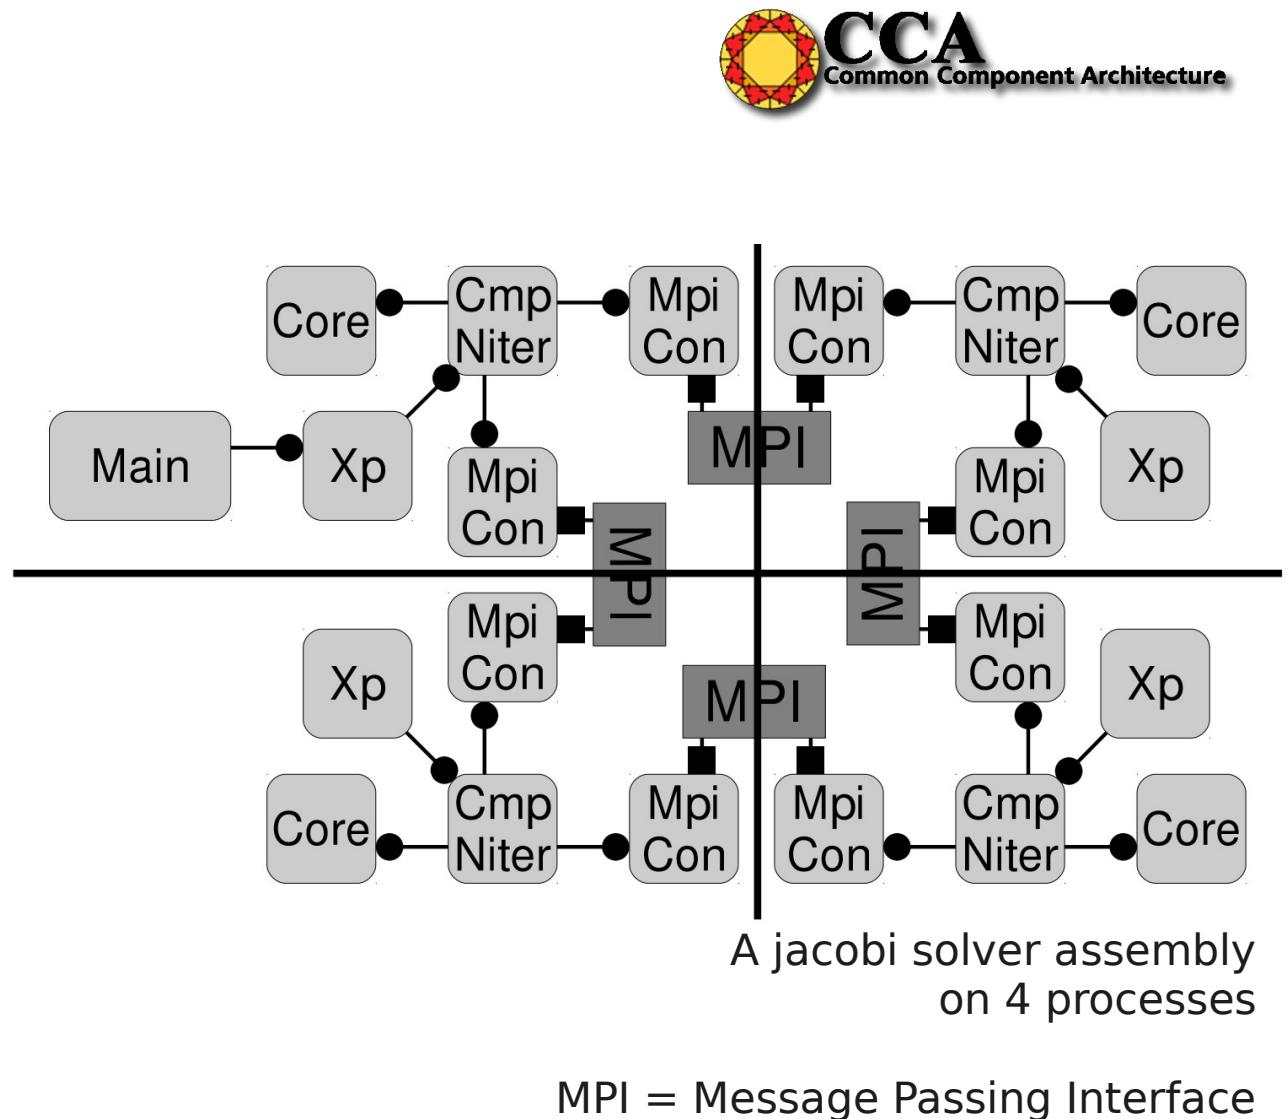
HPC Component Models

Examples:

- CCA
- L2C (Avalon team)

Typically:

- low-level
 - C++/FORTRAN-level abstractions
 - non-hierarchical
- distributed
 - eg, message passing, remote method call
 - process abstraction



Problem: Dynamic HPC Applications

Applications with...

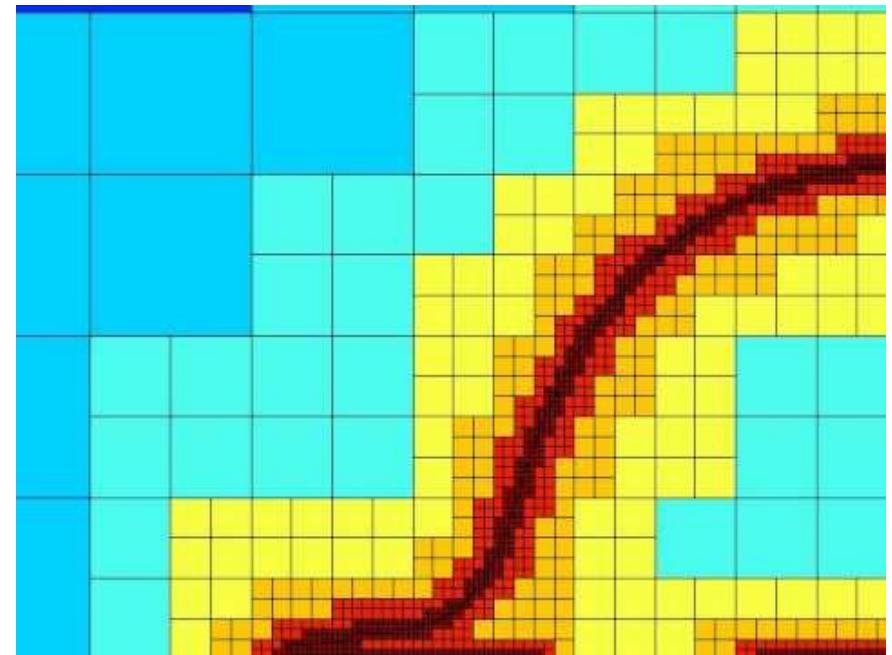
- dynamic communication topology
- dynamic data structure

Examples

- Adaptive Mesh Refinement (AMR)
- dynamic load balancing

Not supported by HPC component models

- reconfiguration needed



AMR mesh,
varying resolution

Goal of this talk:
HPC reconfigurable component model
reconfigurable
scalable (distributed, concurrent reconfiguration)
good SE properties (reuse, separation of concerns)

Plan of the Talk

Context and related works

- Related work
- Our proposition

Presentation of the model, DirectMOD

- Model elements
- Putting it all together

Implementation: DirectL2C

- Code
- Performance

Extensions of DirectMOD

- Efficient locking
- High-level language(s) for transformations

Conclusions and perspectives

Plan of the Talk

Context and related works

- Related work
- Our proposition

Presentation of the model, DirectMOD

- Model elements
- Putting it all together

Implementation: DirectL2C

- Code
- Performance

Extensions of DirectMOD

- Efficient locking
- High-level language(s) for transformations

Conclusions and perspectives

Reconfigurable Component Models

From the literature:

	Examples	Locking and representation	Scalable?	Reconf SE properties
No reconfiguration support	CCA, L2C	none	up to the user	poor
Global reconfiguration	global MAPE	global	no	good
Composite-level controllers	Fractal, SOFA	composite-level	sometimes	sometimes

Important parameters

- locking granularity
 - scalability
- representation granularity
 - ease of use

No model provides both

- scalable approach
- good SE properties for reconfigurable assemblies
 - reuse
 - separation of concerns

Our proposal: main ideas

Let users define locking units (*domains*)

- custom granularity / distribution → performance

Separate locking units from transformations

- transformation on standalone representation
- assembly-level explicit mapping to domain contents
(using *transformation adapters*)

	Locking units	Representation for transformations	Scalable?	Reconf SE properties
Domain-based	user-defined		yes	
Transformation adapters		standalone		good

Plan of the Talk

Context and related works

- Related work
- Our proposition

Presentation of the model, DirectMOD

- Model elements
- Putting it all together

Implementation: DirectL2C

- Code
- Performance

Extensions of DirectMOD

- Efficient locking
- High-level language(s) for transformations

Conclusions and perspectives

A Formal Model

DirectMOD : formal model

- full syntax
- transformation semantics

Benefits

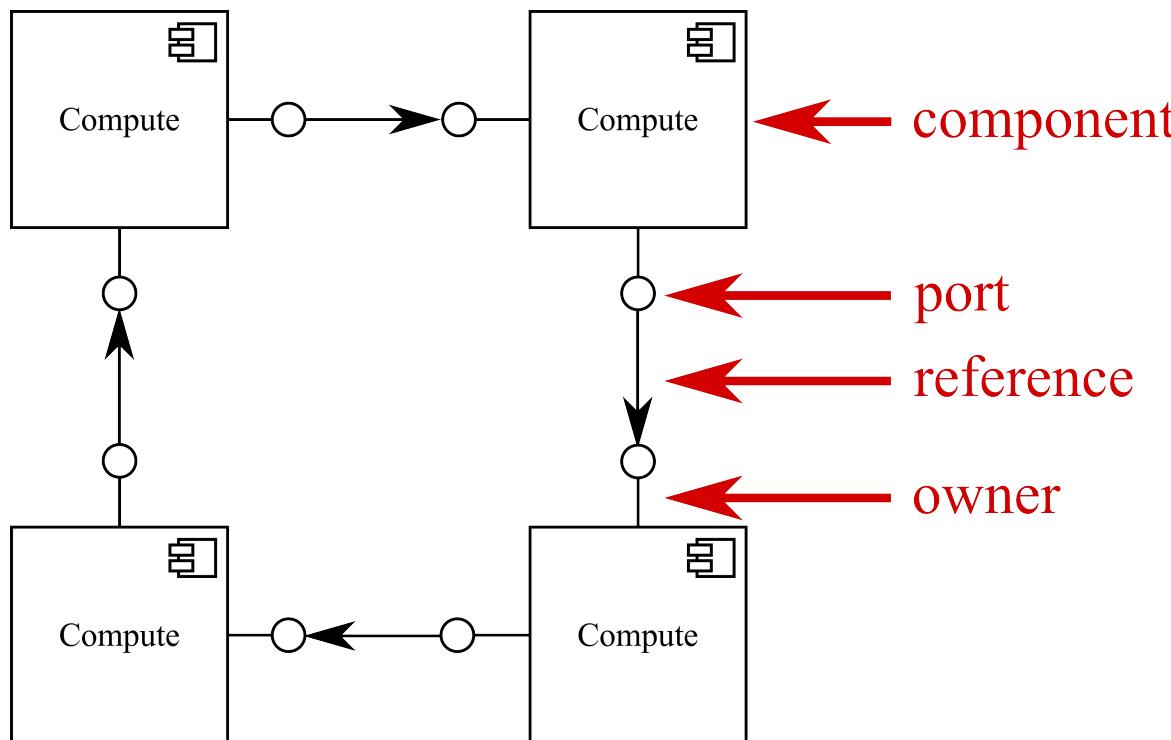
- unambiguous specification
- tech-agnostic
- runtime representation

In addition :

- resource model
(see paper)
- call stack operational
semantics (see paper)



DirectMOD Assembly Model



Elements

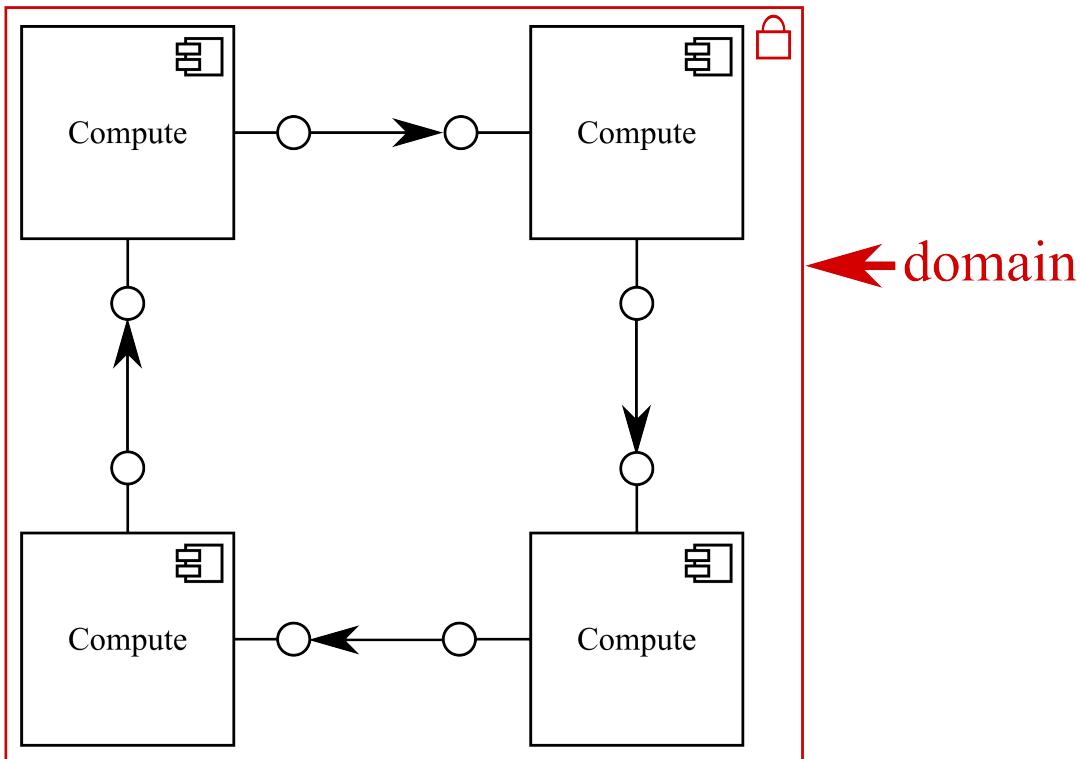
- components
- ports

Relations

- point-to-point references
- owner (component-port relation)

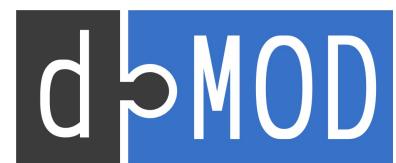


DirectMOD Domains

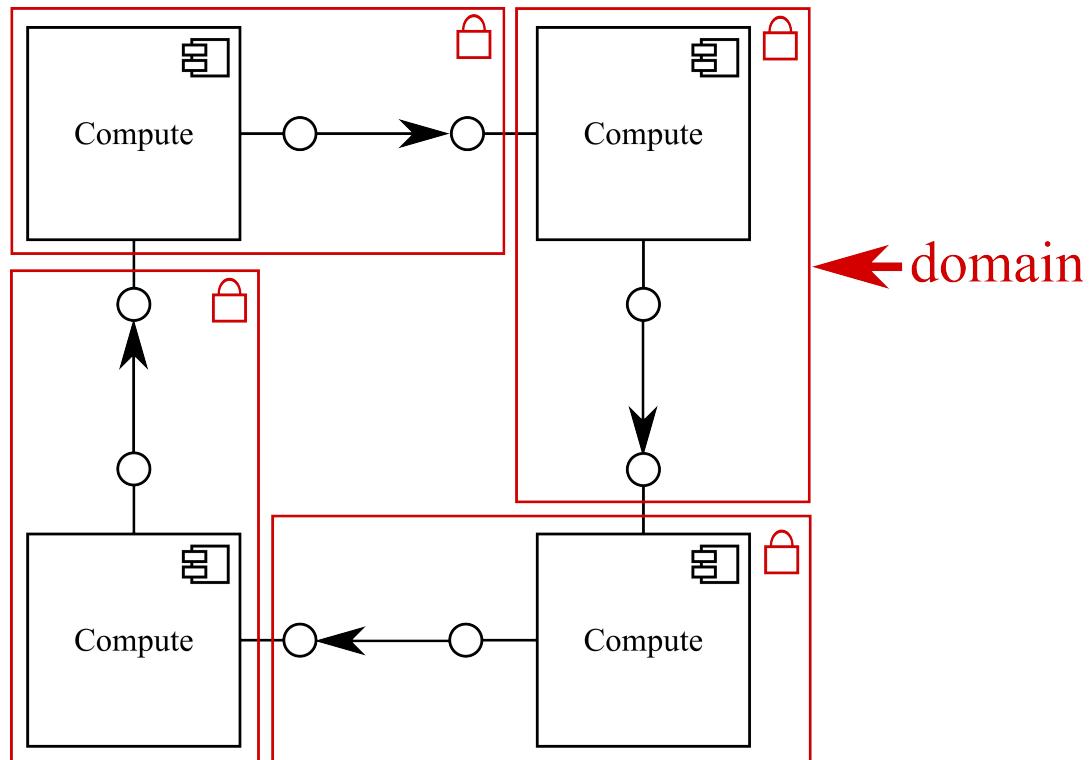


New element: domains

- manage a subassembly
- unit of **locking**
- unit of internal **representation**
- reconfigure their contents

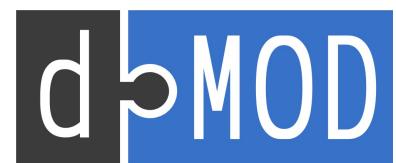


DirectMOD Domains

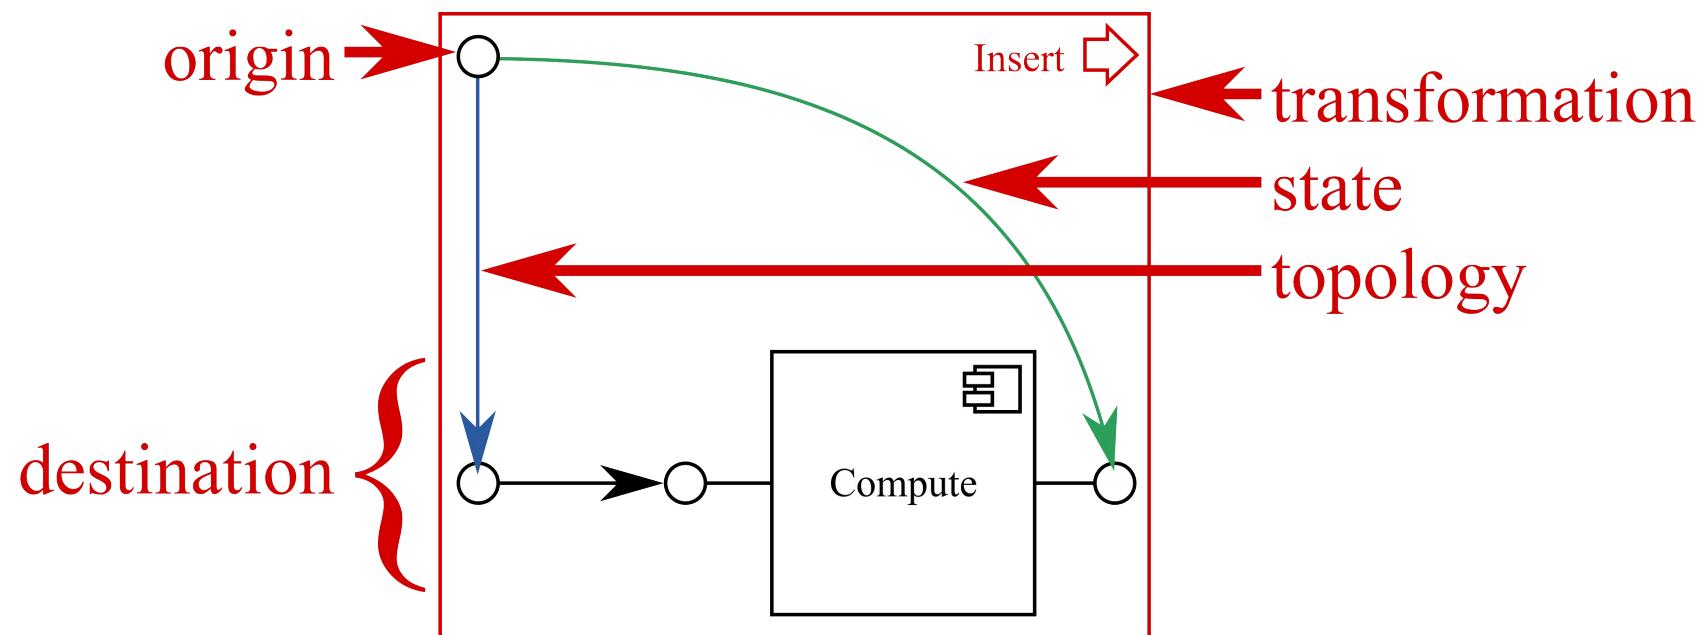


New element: domains

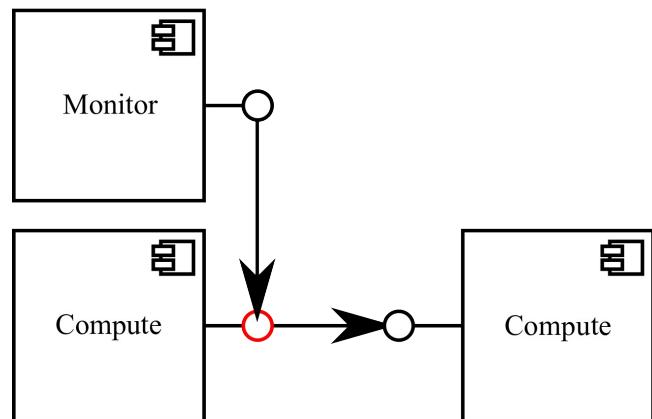
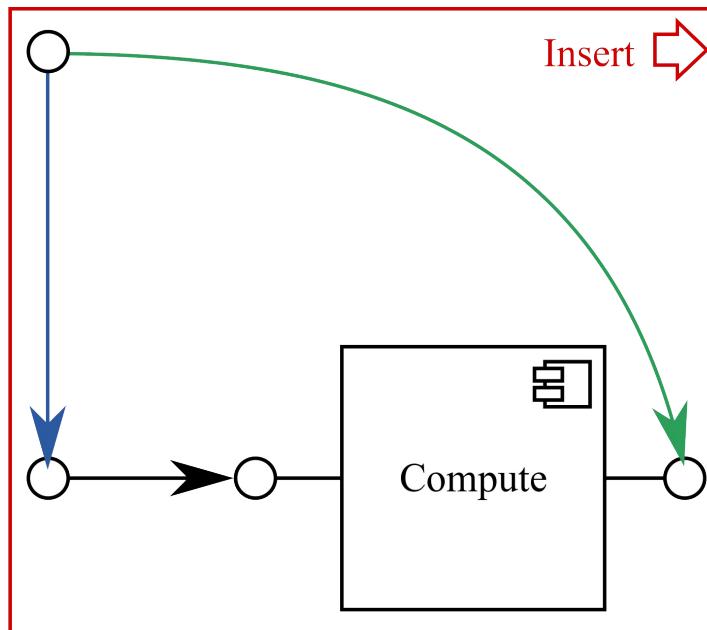
- manage a subassembly
- unit of **locking**
- unit of internal **representation**
- reconfigure their contents



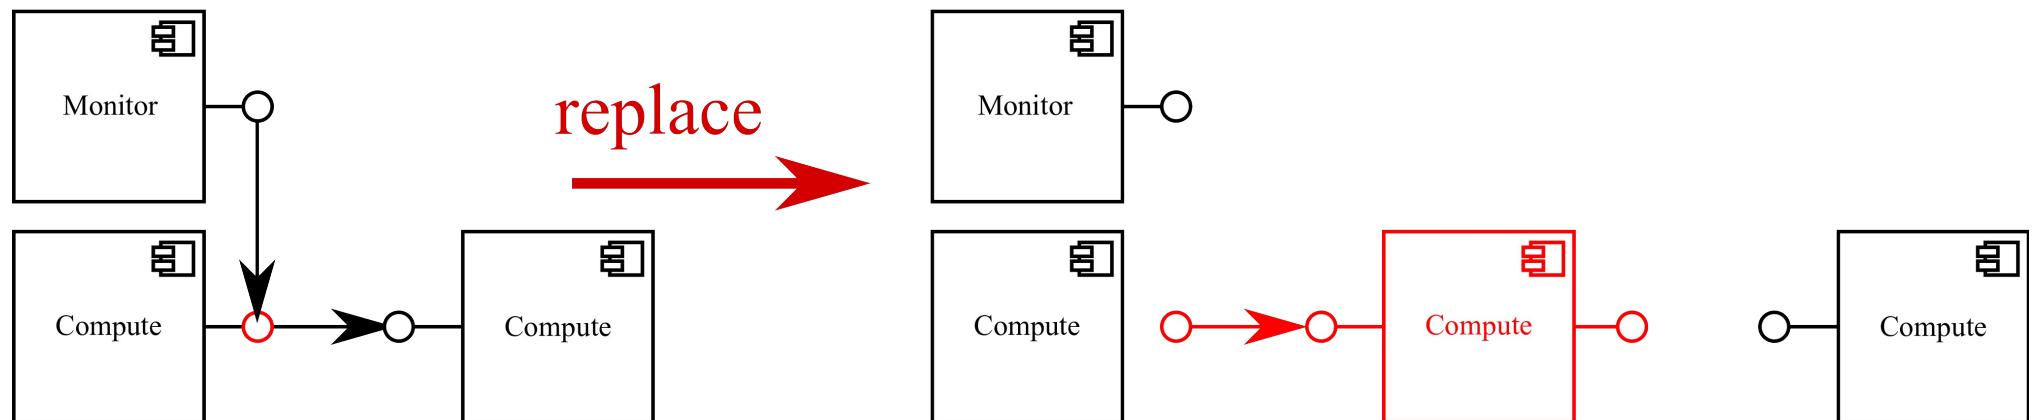
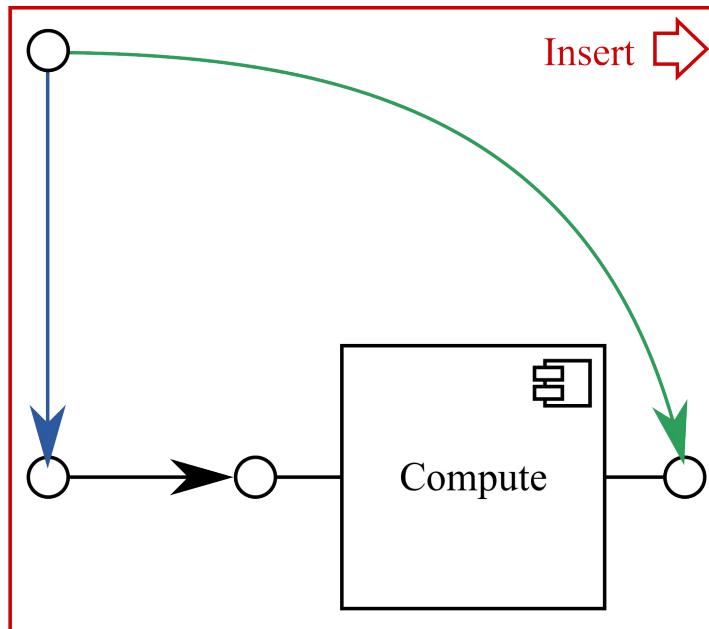
DirectMOD Transformations



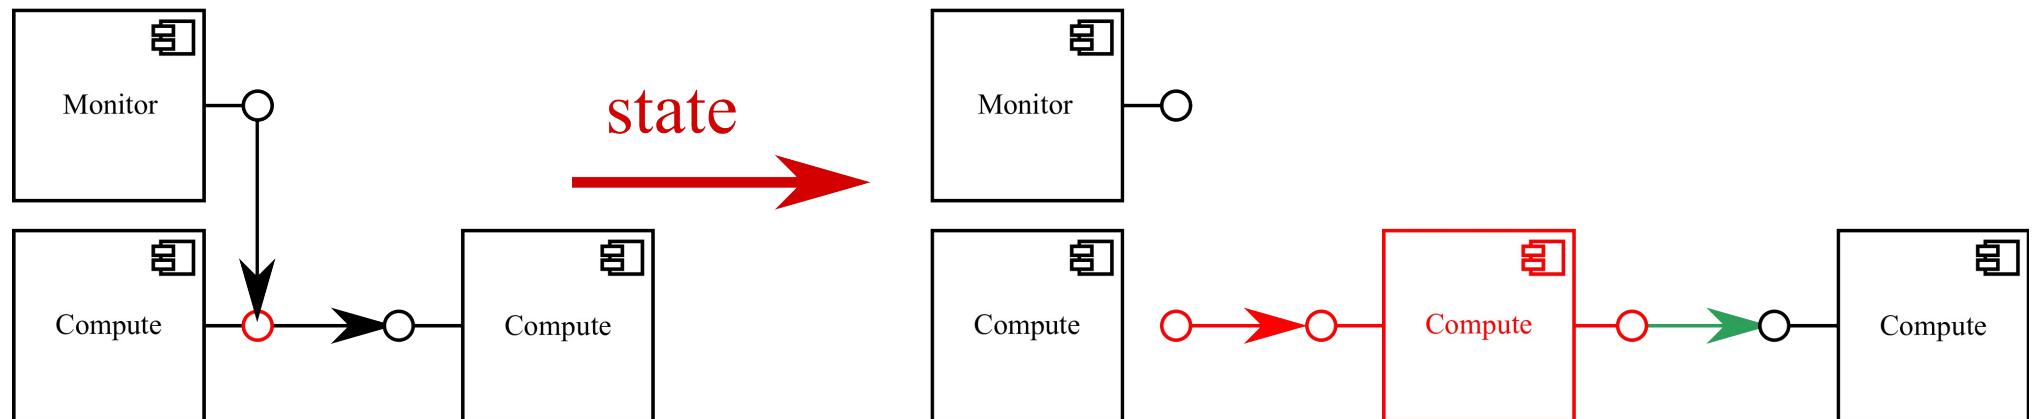
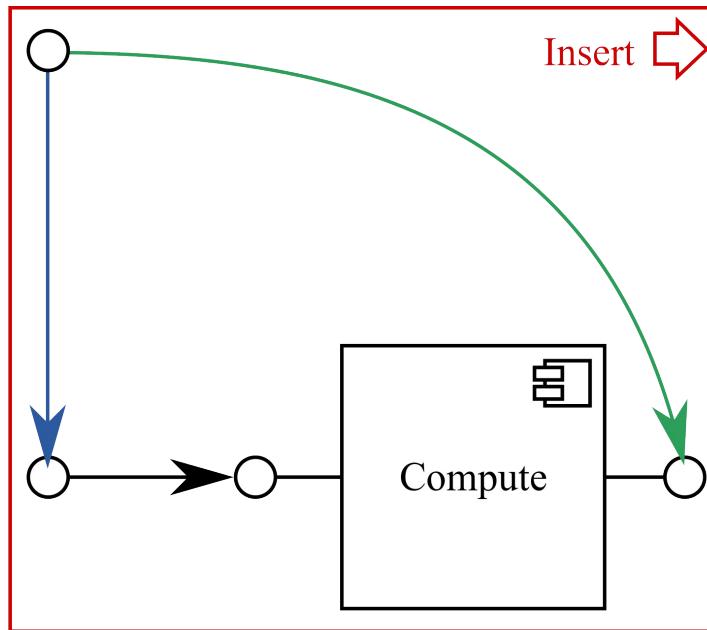
DirectMOD Transformations



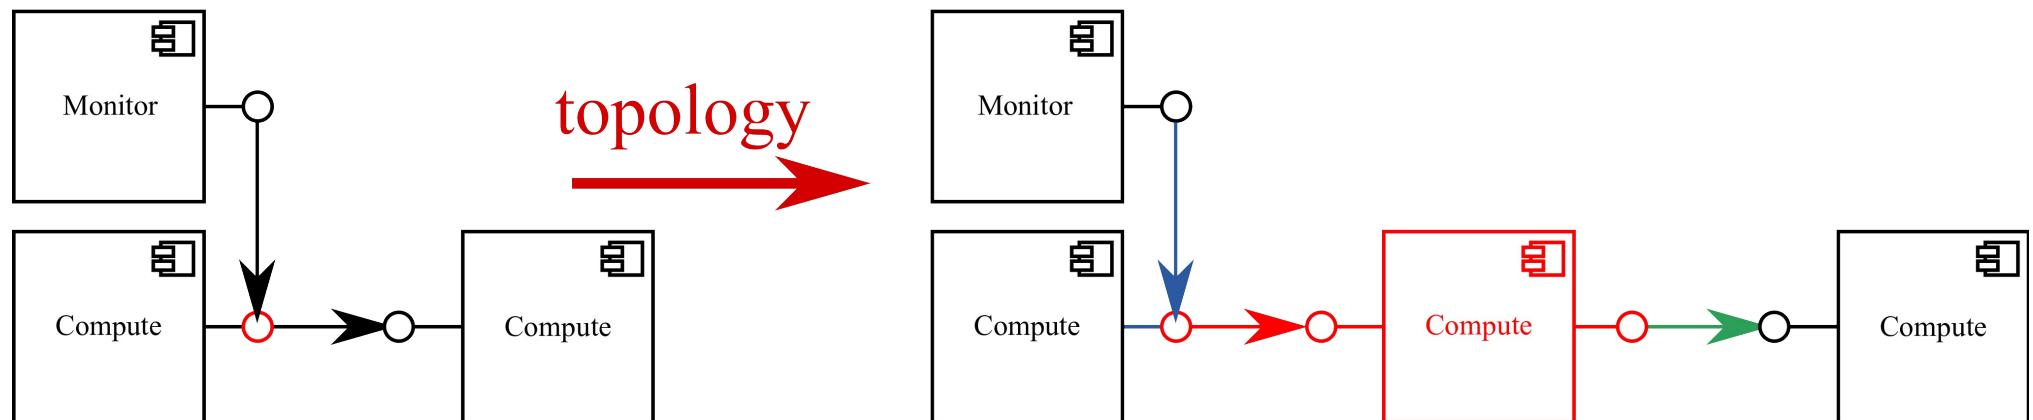
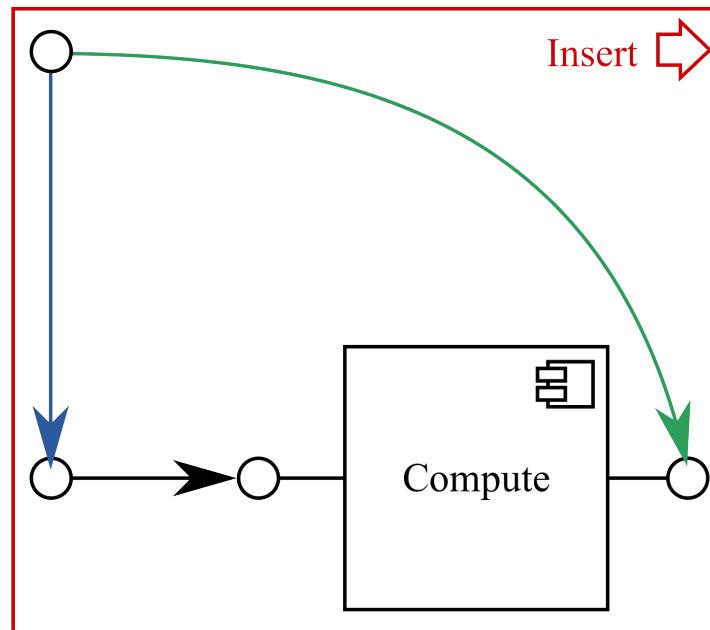
DirectMOD Transformations



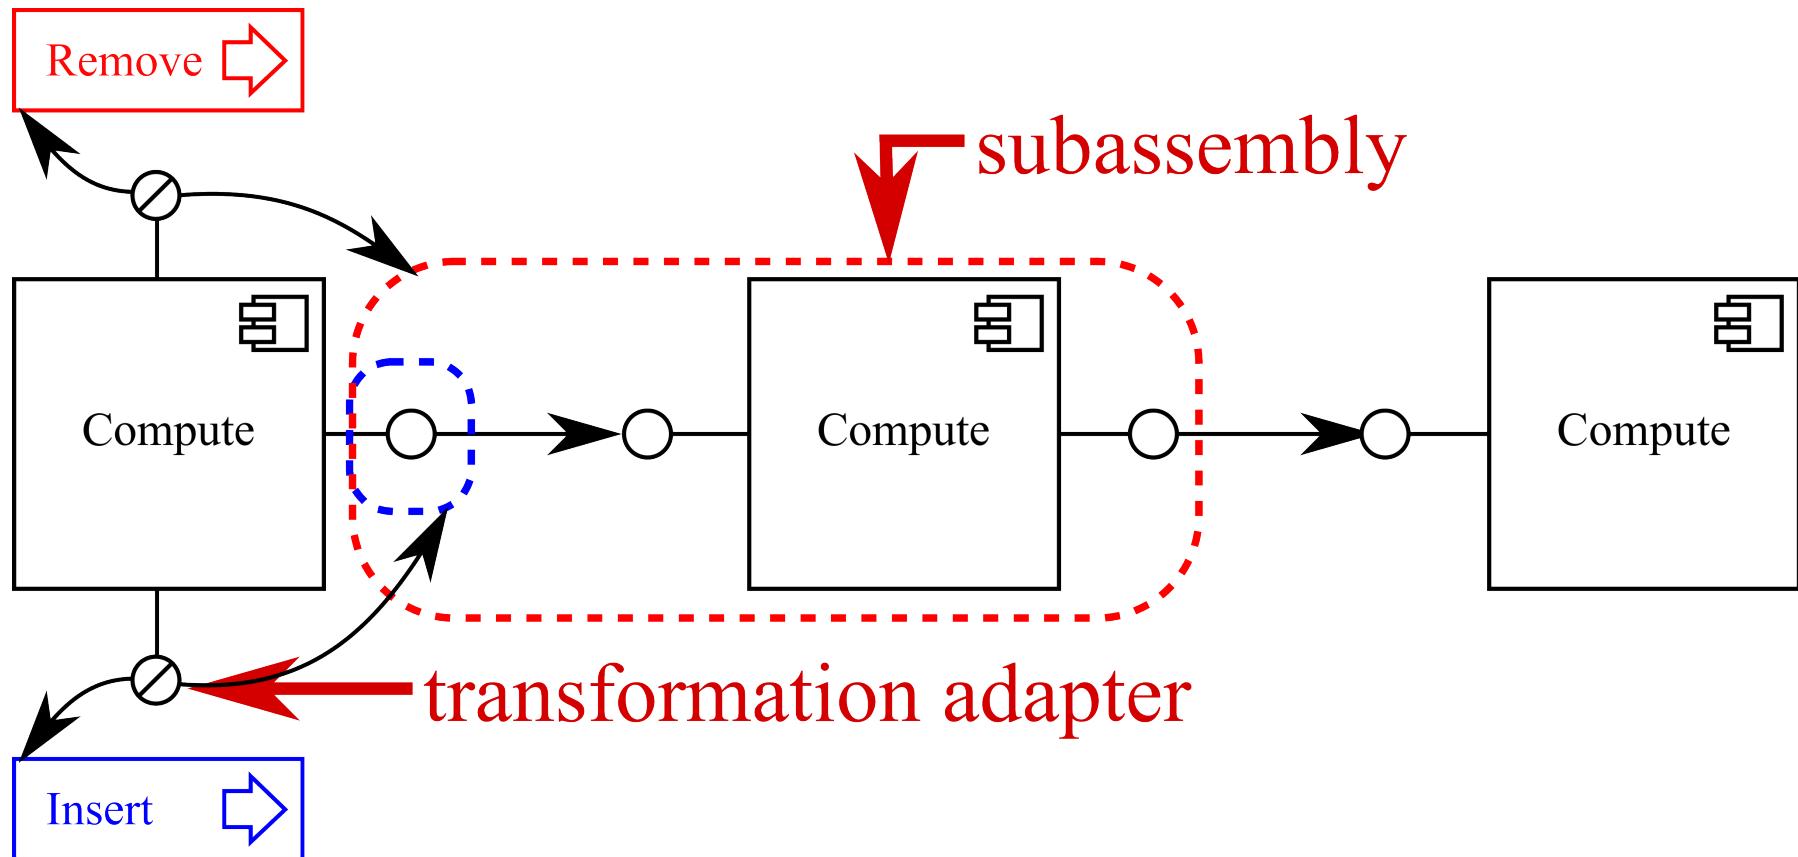
DirectMOD Transformations



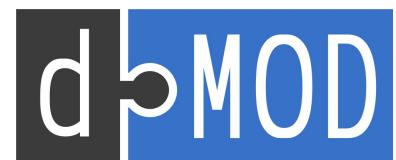
DirectMOD Transformations



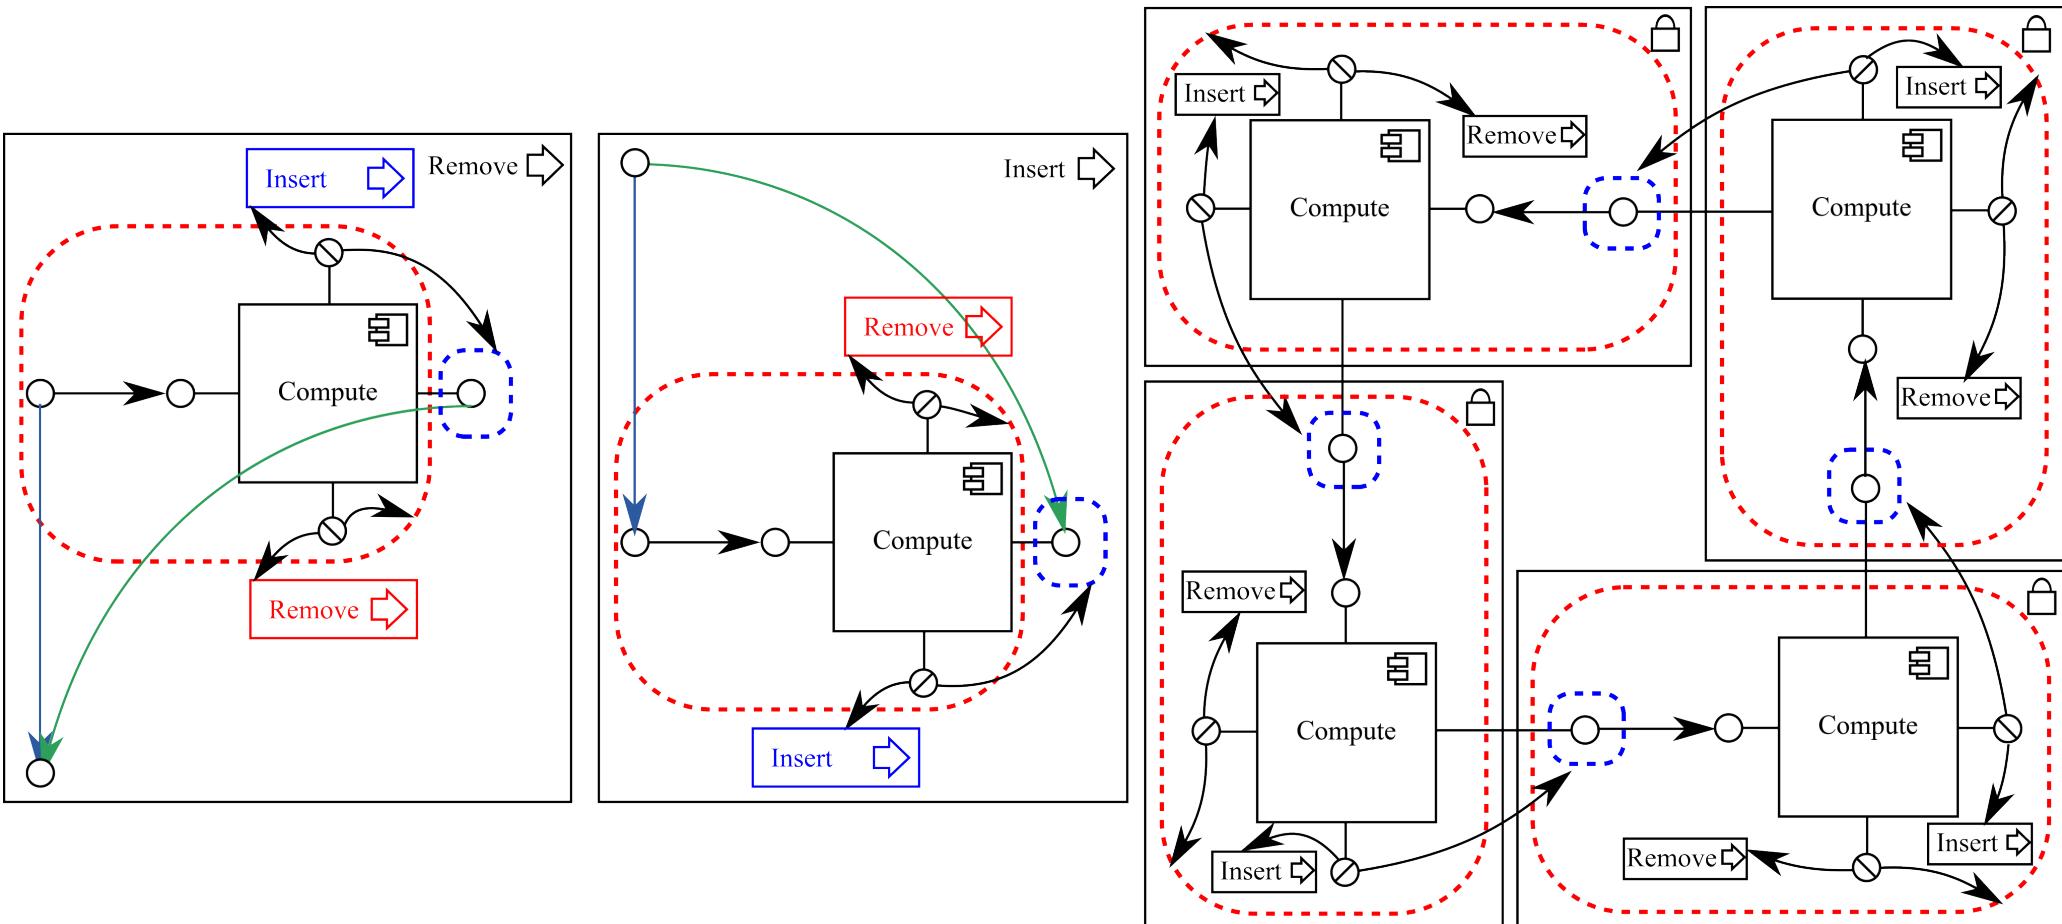
DirectMOD Transformation Adapters



- special kind of port
- reference to transformation and application subassembly

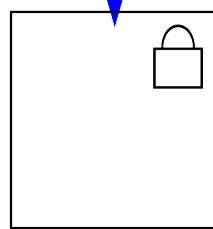
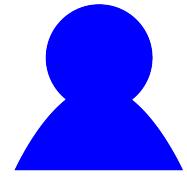


Putting Everything Together



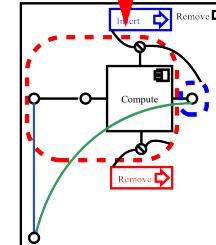
DirectMOD Programming Model

Locking/synchro
specialist



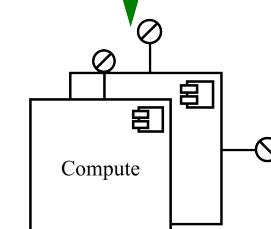
Locking
algorithms

Transformation
programmers



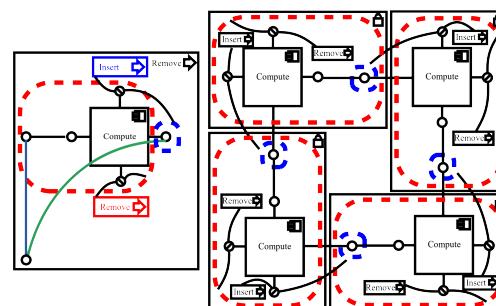
Assembly
transformations

Component
programmers



Components

End user



Reconfigurable
assembly



Plan of the Talk

Context and related works

- Related work
- Our proposition

Presentation of the model, DirectMOD

- Model elements
- Putting it all together

Implementation: DirectL2C

- Code
- Performance

Extensions of DirectMOD

- Efficient locking
- High-level language(s) for transformations

Conclusions and perspectives

A C++/MPI Implementation

DirectL2C

- DirectMOD implementation
- extension of L2C
- uses traditional HPC tech
 - C++
 - MPI (Message Passing Interface)
 - threads

	Components	Files	mLOC
L2C	37	79	4570
DirectL2C	3	10	1118

L2C provides

- C++ components with zero overhead
- basic component operations
- assembly deployment
 - distributed

DirectL2C provides

- transformation parsing and execution
- interface and locking APIs



Evaluation: Easy to Write?

Implemented: ring assembly

- insert new component
- remove component

Preliminary implementation

- not yet fully optimized
- shortest possible code

Ring assembly LOC

Function	C++ LOC
Transformation	8
Non-functional sync	20
Code instrumentation	13
L2C overhead	7
DirectL2C overhead	6
Functional code	31
Other	3
TOTAL	88

- short and easy transformation code
- complex synchronization code

Transformation Pseudocode

Inputdata: portName, reconfPortName, resourceName

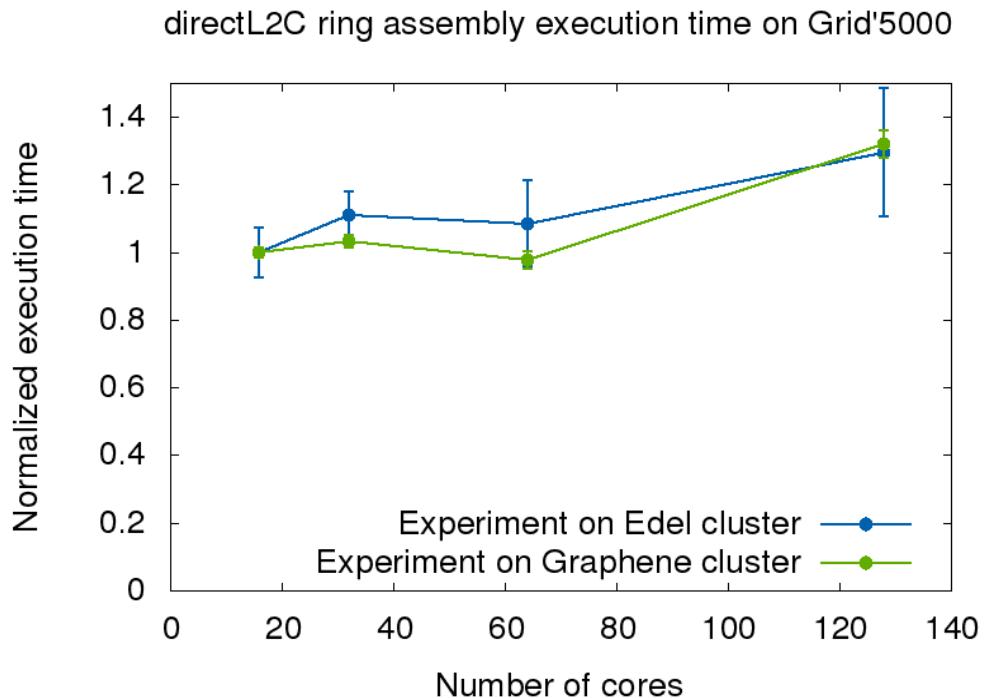
```
Direct::transformation insert();  
insert.create("Cp", "newCp", resourceName);  
//ports of created component have implicit names  
insert.connect(portName, "newCpLeft");  
insert.statePort(portName, "newCpRight");  
insert.topoPort(reconfPortName, "newCpReconf");
```

Evaluation: scalable?

Testing our ring assembly

- one component per core at startup
- fixed number of *insert* and *remove* transformations per starting component
- one domain per component (fully distributed)

Preliminary experiments



- on Grid'5000 clusters
- acceptable scalability up to 128 cores

Plan of the Talk

Context and related works

- Related work
- Our proposition

Presentation of the model, DirectMOD

- Model elements
- Putting it all together

Implementation: DirectL2C

- Code
- Performance

Extensions of DirectMOD

- Efficient locking
- High-level language(s) for transformations

Conclusions and perspectives

Some questions left unanswered

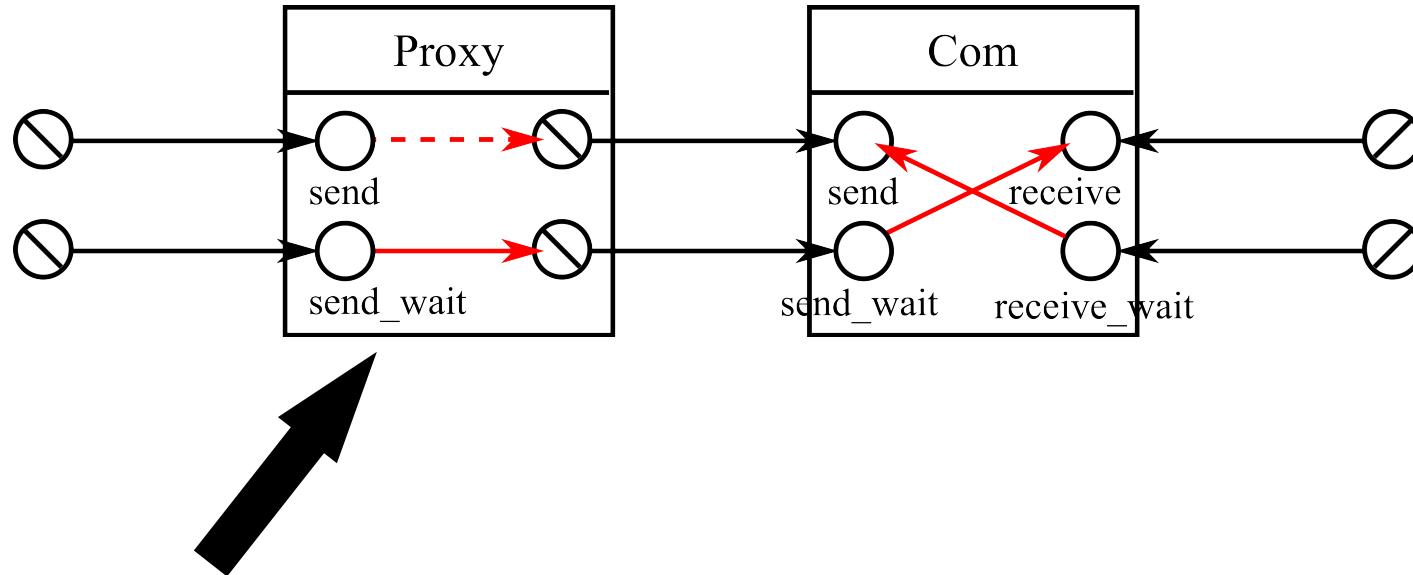
■ what about locking domains?

- looks difficult
- poor reuse

■ what about transformations?

- how to write them? language?
- what about genericity?
- ... hierarchy?

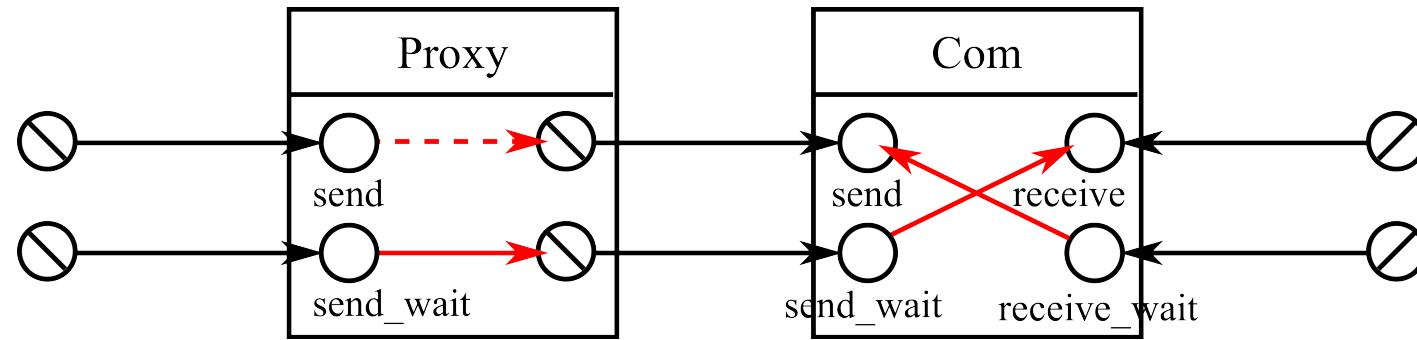
Stopping Components



**This assembly is running...
... how to stop it?**

- no control info
- performance?
- deadlocks?
- call stacks

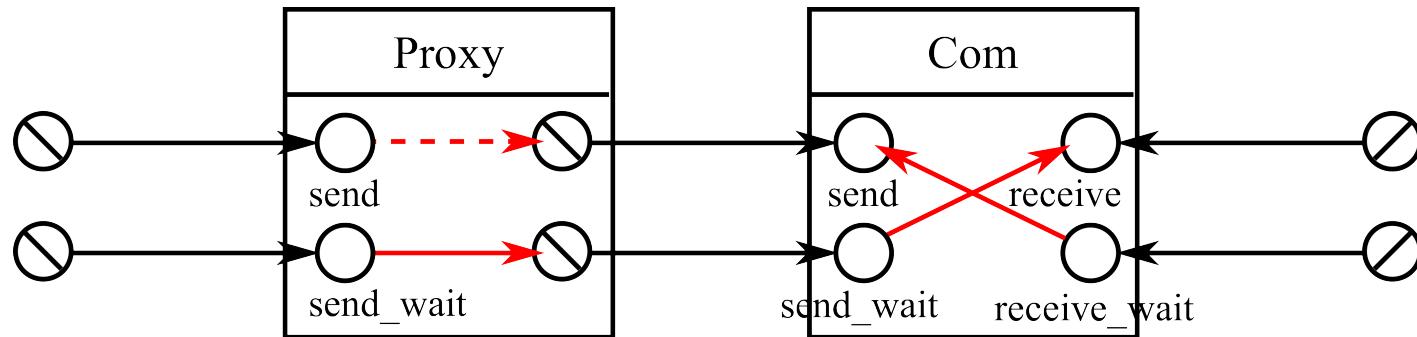
Locking use ports



Approach: locking use ports

- no more calls once locked
- mutex-like behaviour
- in what order?
- deadlocks?

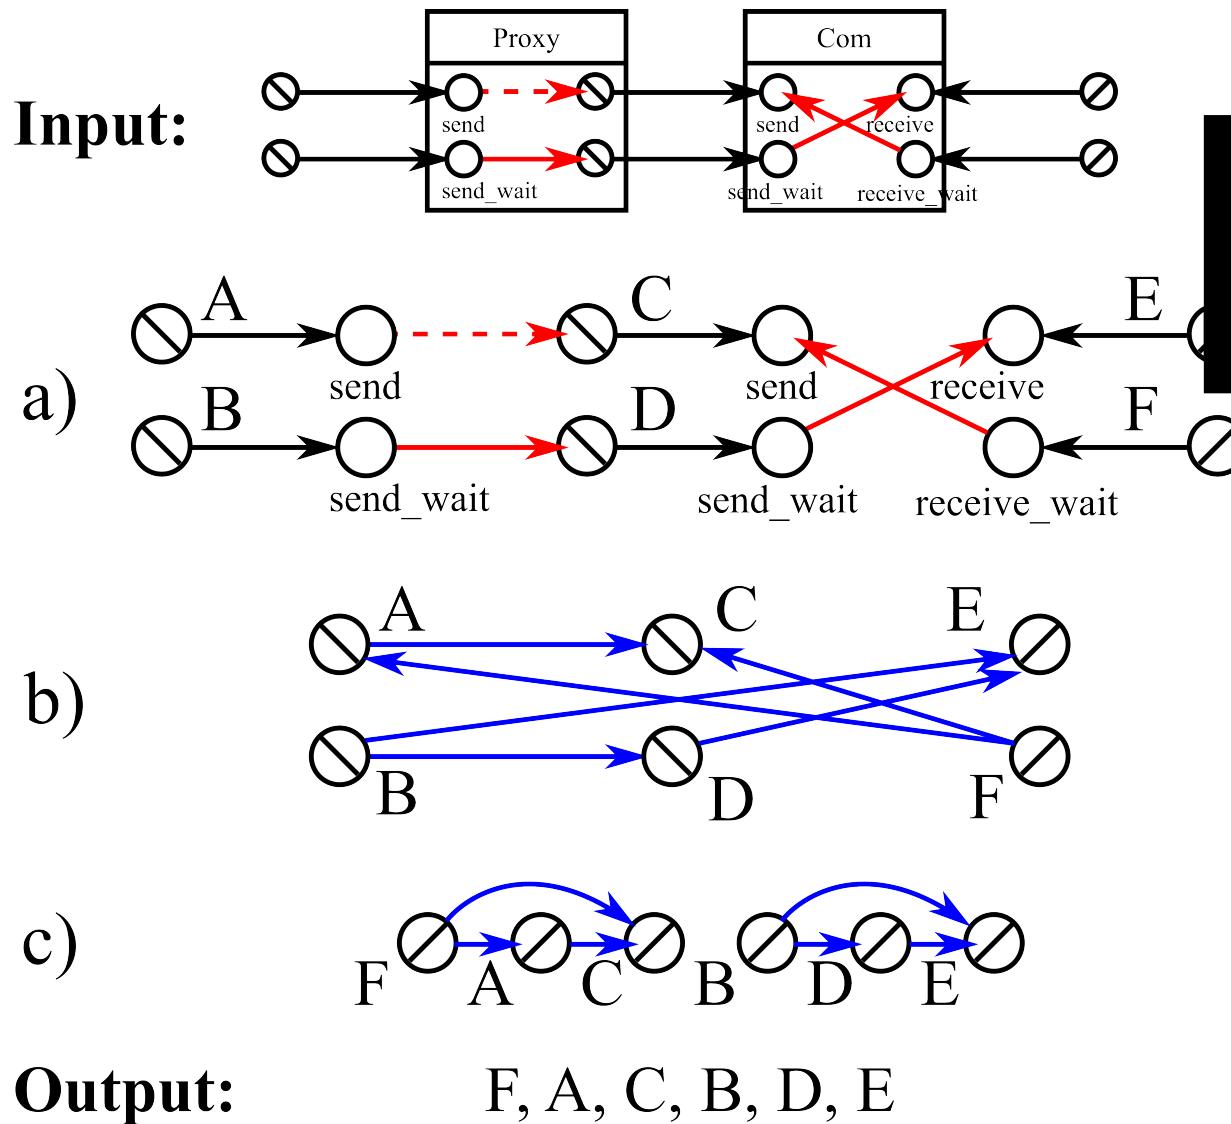
Adding metadata



First step: add metadata

- minimal
- call dependencies
- termination dependencies

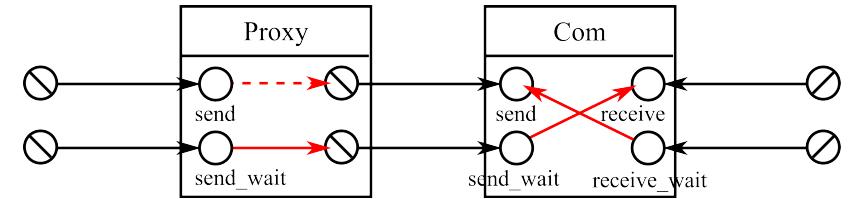
Second Step: Locking Order Algorithm



It Works for Stencil Applications

■ stencil applications

- simple control
- concurrent locking of connections

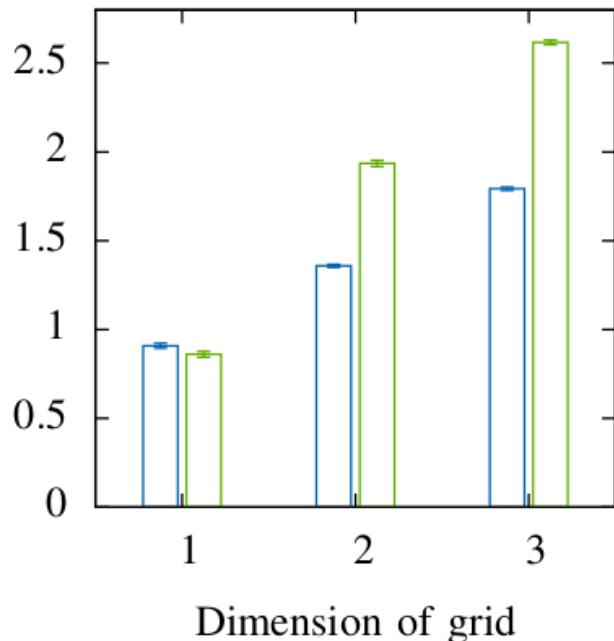


■ useful automation: connection variants

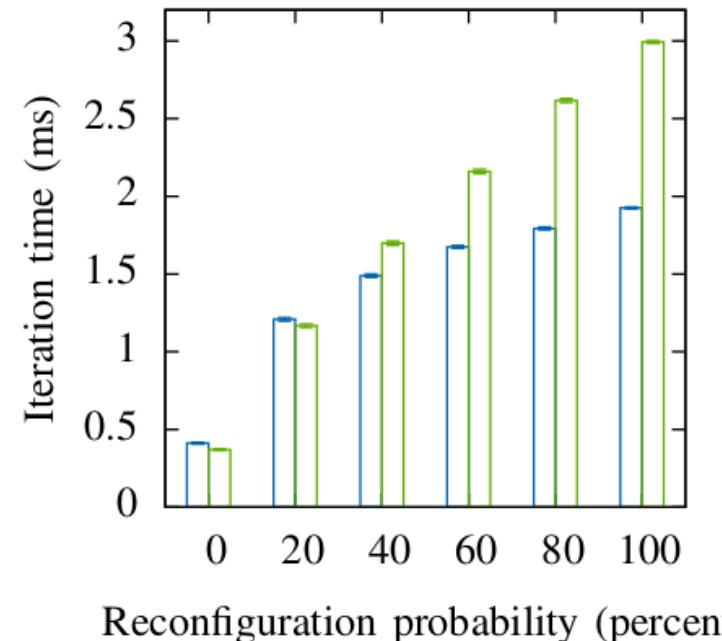
- local, distributed
- proxies
- 1-to-n connections

Locking Performance

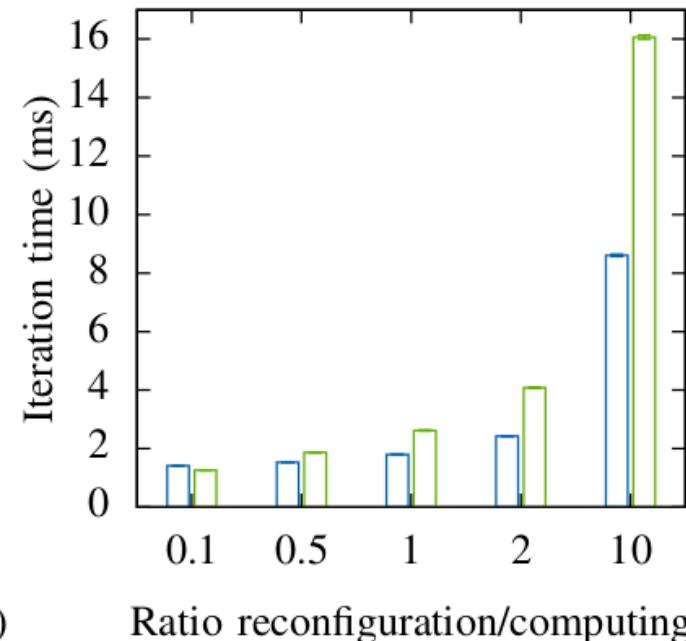
Locking performance



Locking performance



Locking performance



- 1/2/3D grid benchmark w/ locking only
- versus component-wise approach

High-level transformation language(s)

■ High-level languages

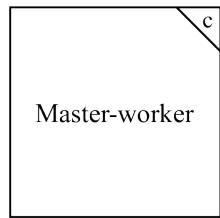
- hierarchy
- genericity
- decent compact language for humans (ie, not XML)

■ Transformation to low-level

- off-line process
- optimization during transformation

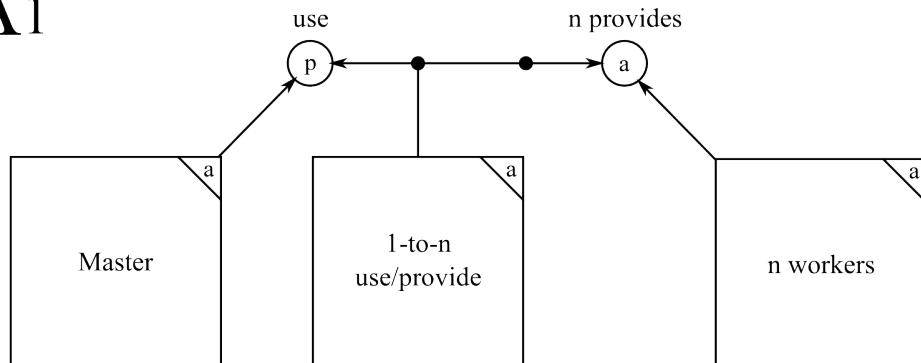
High-level transformation language(s)

A0



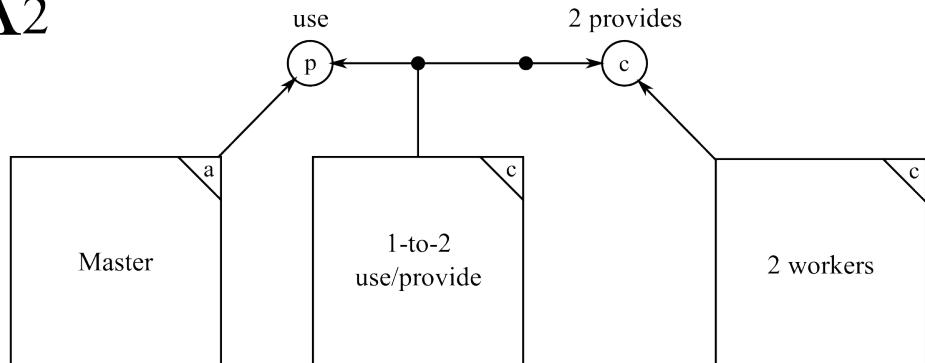
$V(A_0, \Sigma) = \{$
 $\text{im}_c(\text{Master-Worker})\}$

A1



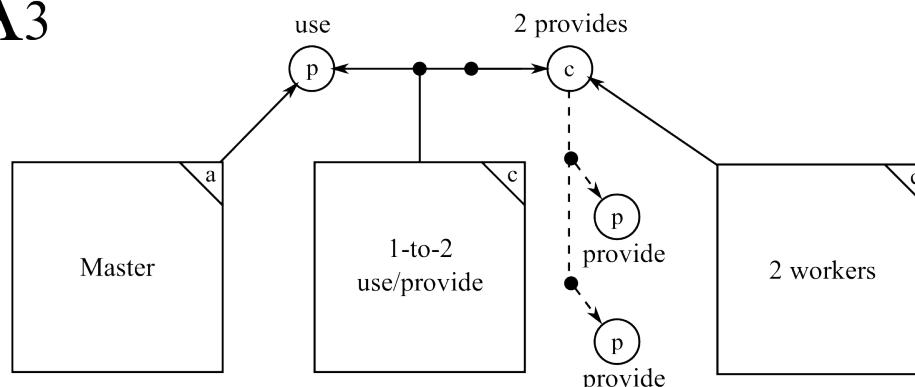
$V(A_1, \Sigma) = \{$
 $\text{sp}_c(\text{Master}, \text{Master A})$
 $\text{sp}_c(\text{Master}, \text{Master B})$
 $\text{sp}_c(n \text{ workers}, 2 \text{ workers})$
 $\text{sp}_c(n \text{ workers}, 1 \text{ worker})$
 $\text{sp}_c(n \text{ workers}, 2 \text{ identical workers})$
 $\text{sp}_c(n \text{ provides}, 2 \text{ provides})$
 $\text{sp}_c(n \text{ provides}, 1 \text{ provide})$
 $\text{sp}_c(1\text{-to-}n \text{ use/provide}, 1\text{-to-}2 \text{ use/provide})$
 $\text{sp}_c(1\text{-to-}n \text{ use/provide}, \text{use/provide})$
\}

A2



$V(A_2, \Sigma) = \{$
 $\text{sp}_c(\text{Master}, \text{Master A})$
 $\text{sp}_c(\text{Master}, \text{Master B})$
 $\text{im}_c(2 \text{ provides})$
\}

A3

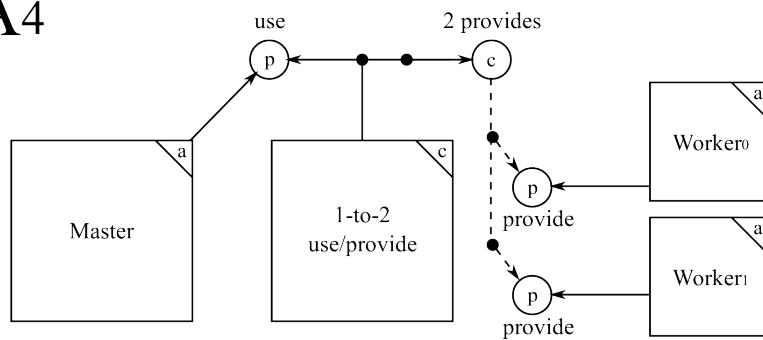


$V(A_3, \Sigma) = \{$
 $\text{sp}_c(\text{Master}, \text{Master A})$
 $\text{sp}_c(\text{Master}, \text{Master B})$
 $\text{im}_c(2 \text{ workers})$
 $\text{im}_c(1\text{-to-}2 \text{ use/provide})$
\}

$V(A_5, \Sigma) = \{$
 $\text{sp}_c(\text{Master}, \text{Master A})$
 $\text{sp}_c(\text{Master}, \text{Master B})$
 $\text{sp}_c(\text{Worker}_0, \text{Worker A})$
 $\text{sp}_c(\text{Worker}_1, \text{Worker A})$
 $\text{sp}_c(\text{Worker}_0, \text{Worker B})$
 $\text{sp}_c(\text{Worker}_1, \text{Worker B})$
\}

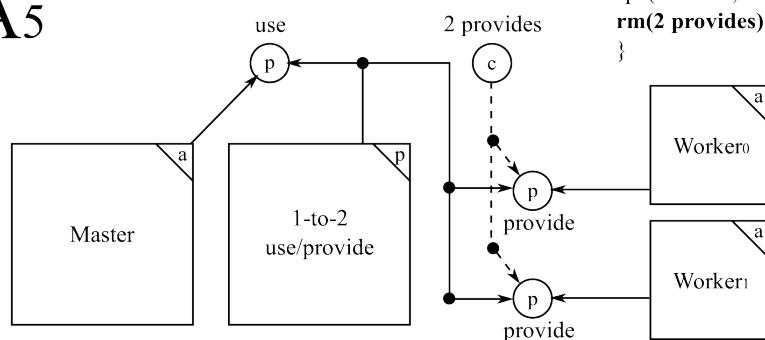
High-level transformation language(s)

A4



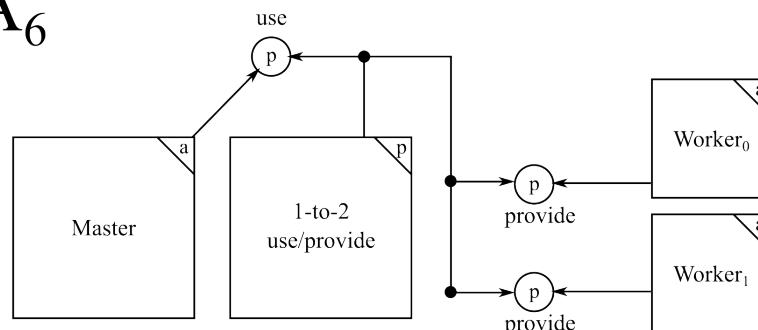
$V(A_4, \Sigma) = \{$
 $\text{sp}_c(\text{Master}, \text{Master A})$
 $\text{sp}_c(\text{Master}, \text{Master B})$
 $\text{spc}(\text{Worker}_0, \text{Worker A})$
 $\text{spc}(\text{Worker}_1, \text{Worker A})$
 $\text{spc}(\text{Worker}_0, \text{Worker B})$
 $\text{spc}(\text{Worker}_1, \text{Worker B})$
imc(1-to-2 use/provide)
 $\}$

A5



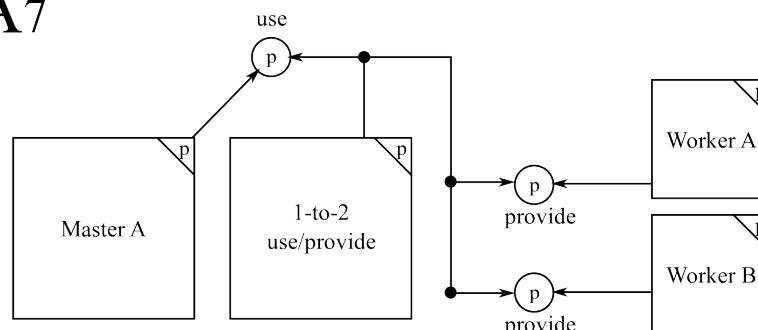
$\text{spc}(\text{Worker}_0, \text{Worker B})$
 $\text{spc}(\text{Worker}_1, \text{Worker B})$
rm(2 provides)
 $\}$

A6



$V(A_6, \Sigma) = \{$
spc(Master, Master A)
 $\text{sp}_c(\text{Master}, \text{Master B})$
spc(Worker}_0, \text{Worker A)}
 $\text{spc}(\text{Worker}_1, \text{Worker A})$
 $\text{spc}(\text{Worker}_0, \text{Worker B})$
spc(Worker}_1, \text{Worker B})
 $\}$

A7



$V(A_7, \Sigma) = \{$
 $\}$

Plan of the Talk

Context and related works

- Related work
- Our proposition

Presentation of the model, DirectMOD

- Model elements
- Putting it all together

Implementation: DirectL2C

- Code
- Performance

Extensions of DirectMOD

- Efficient locking
- High-level language(s) for transformations

Conclusions and perspectives

Conclusion and Perspectives

Presented DirectMOD

- a reconfigurable component model
- formal
- distributed
- + implementation (DirectL2C)
- + efficient locking for stencil applications

Perspectives

- improve evaluation & implementation
 - ongoing work on a complex benchmark
 - experiments on a large platform (DARI allowance)
 - + work on directL2C
- transformation specification
 - genericity
 - compact language